

~~Machine k~~

Runtime - 1

Data Representation

- size & alignment
- structs - holes.
- SVID - layout of structs unions

bit numbering
big / little endian
host order / network order.

Register usage

- free used ~~beginning~~ regs.
- minimize bookkeeping regs.
- caller save / callee save.

SP → top of runtime stack.

(BP)FP → beginning of current frame
(may be virtual if frame fixed)

dynamic (control) link
static (access) link

volatile variables.

ML

1, 2, 4, 8

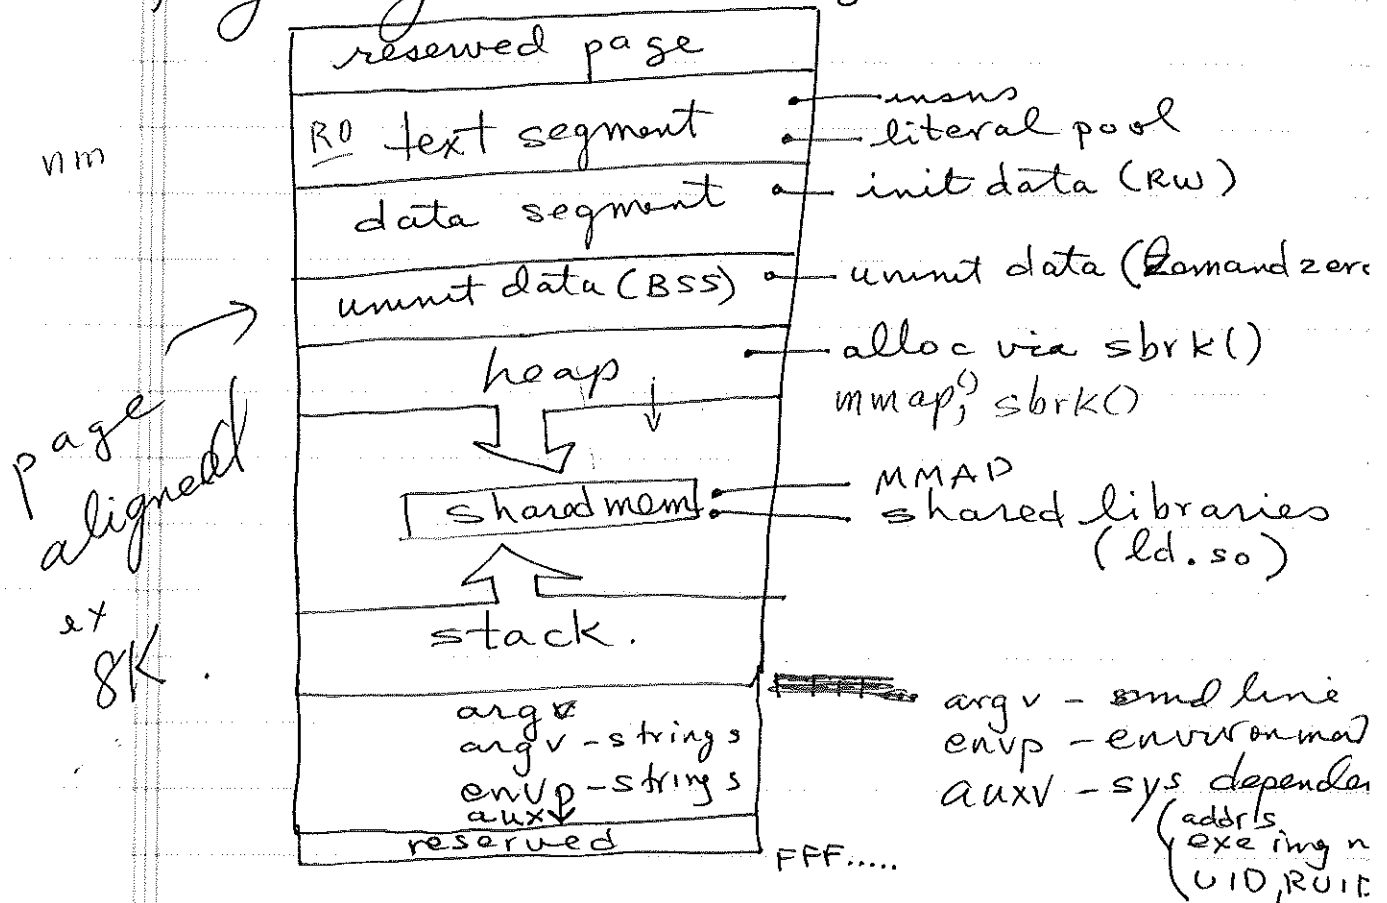
eax rax
ebx
ecx
edx
esi
edi
esp
ebp

rs
r15

~~contents~~

Prog layout

Runtime 2



libraries - *ld static* -
ld dynamic - versioning pr

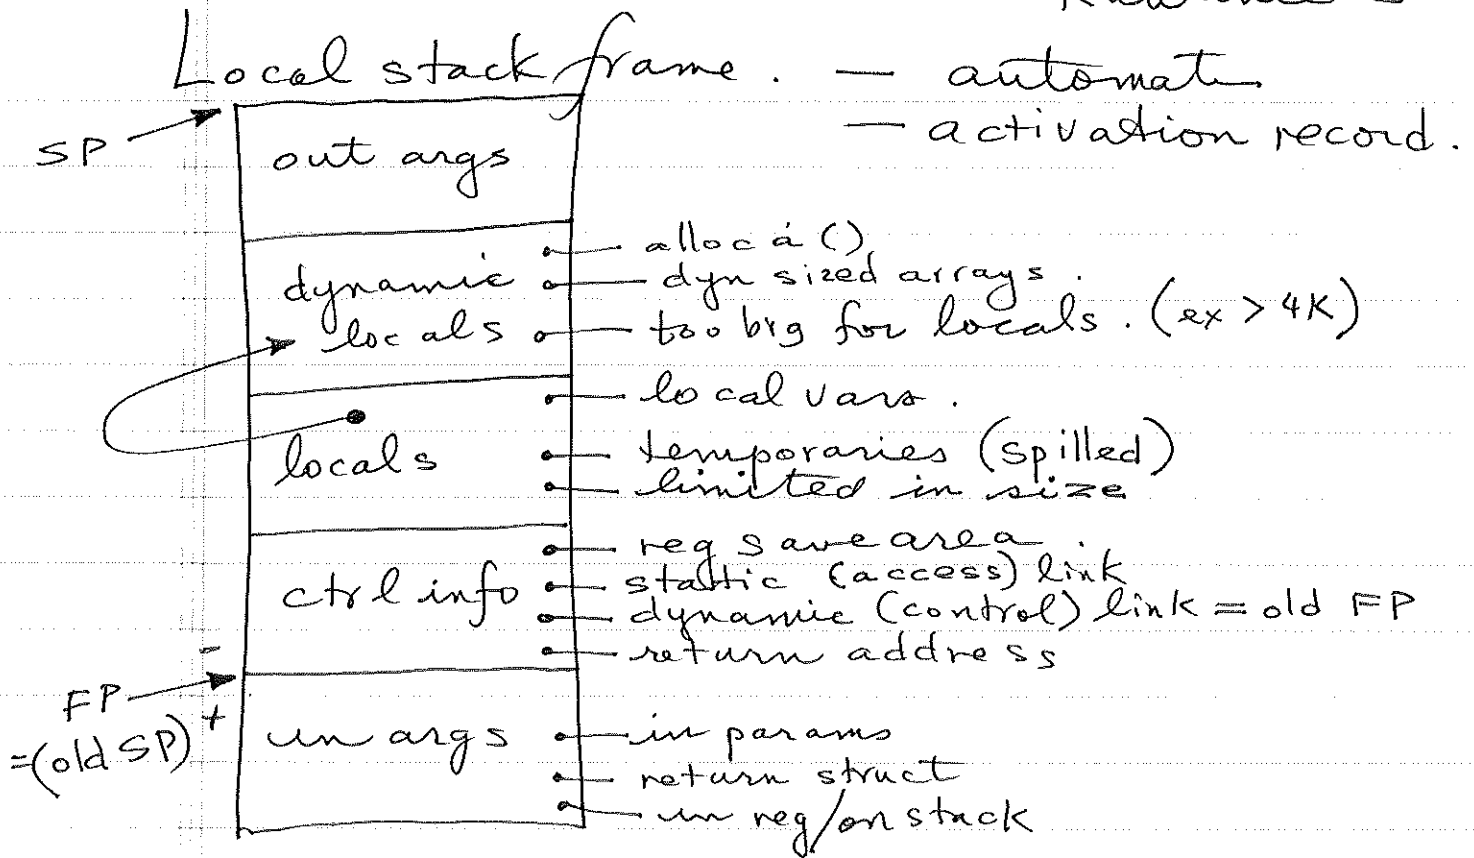
Variables = static | auto | heap.

static - easiest to alloc COMMON BLO
 - addr @ link time
 - problem with recursive fns
 - class - single instance
 - class table static.

scope: fn, file scope, external scope.

extern (vs) static
 C - all fns static.

Runtime 3



recursion requires stack.

control

static link

dynamic link.

global offset table ptr (code sharing)

nested procs:

caller - callee next level


(-1) → static link = caller's frame

(0) → s.l → copy of caller's s.l.

(n > 0) → n static links back

Runtime 4

Calling Seq



call - assemble args.

- regs 00, 01, ..., 05, then stack
- Pentium → reverse push.

- get addr of fn (usu. static)

- ~~store~~ caller save regs if in use

- compute static link

- save ret addr & do call.

prolog - save old frame

- alloc new frame

- save callee save regs if used.

- construct display?

X	X	X	X
---	---	---	---

 ←

epilog - load callee saved regs.

- recover old SP/FP

- set up ret value

- jump to ret addr.

at return - load caller saved regs.


- get return value.

gcc -O6

Runtime 5


Param Passing

- reentrant procs.
- Runtime stack
 - pushed in reverse order
- passed in flat regs. #MIPS regs = 4
 - $p0, p1, p2, p3$ - load regs.
 - rest in some point on stack
 - must be saved
- passed in reg windows
 - #regs on stack = 6
 - windows cycle ~~MIPS = 4~~
- typical - 2 params per fn
 - value of




- kinds
- by value (in) small
 - by result (copy back)
 - by value-result
 - by reference large
 - by name - thunk.
 - by descriptor (ref/ref).
- ~~functions as params.~~

causes
aliasing prob
unless
~~immutable~~
immutable



passing fns

Runtime 6



Access to non-local

local $\Rightarrow -N(FP)$
or $+N(SP)$

global \Rightarrow absolute addr
or: $\pm N(gp)$ global ptr
(per module?)
(per class?)

other: followⁿ static links

display f {

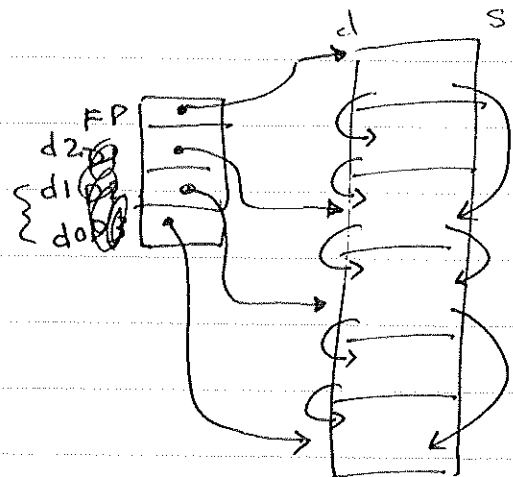
g {

h {

}

}

}



use d0 - most global
d1 - next

- save/restore @ each call.
- how many regs for display?

Static Lexical Scope

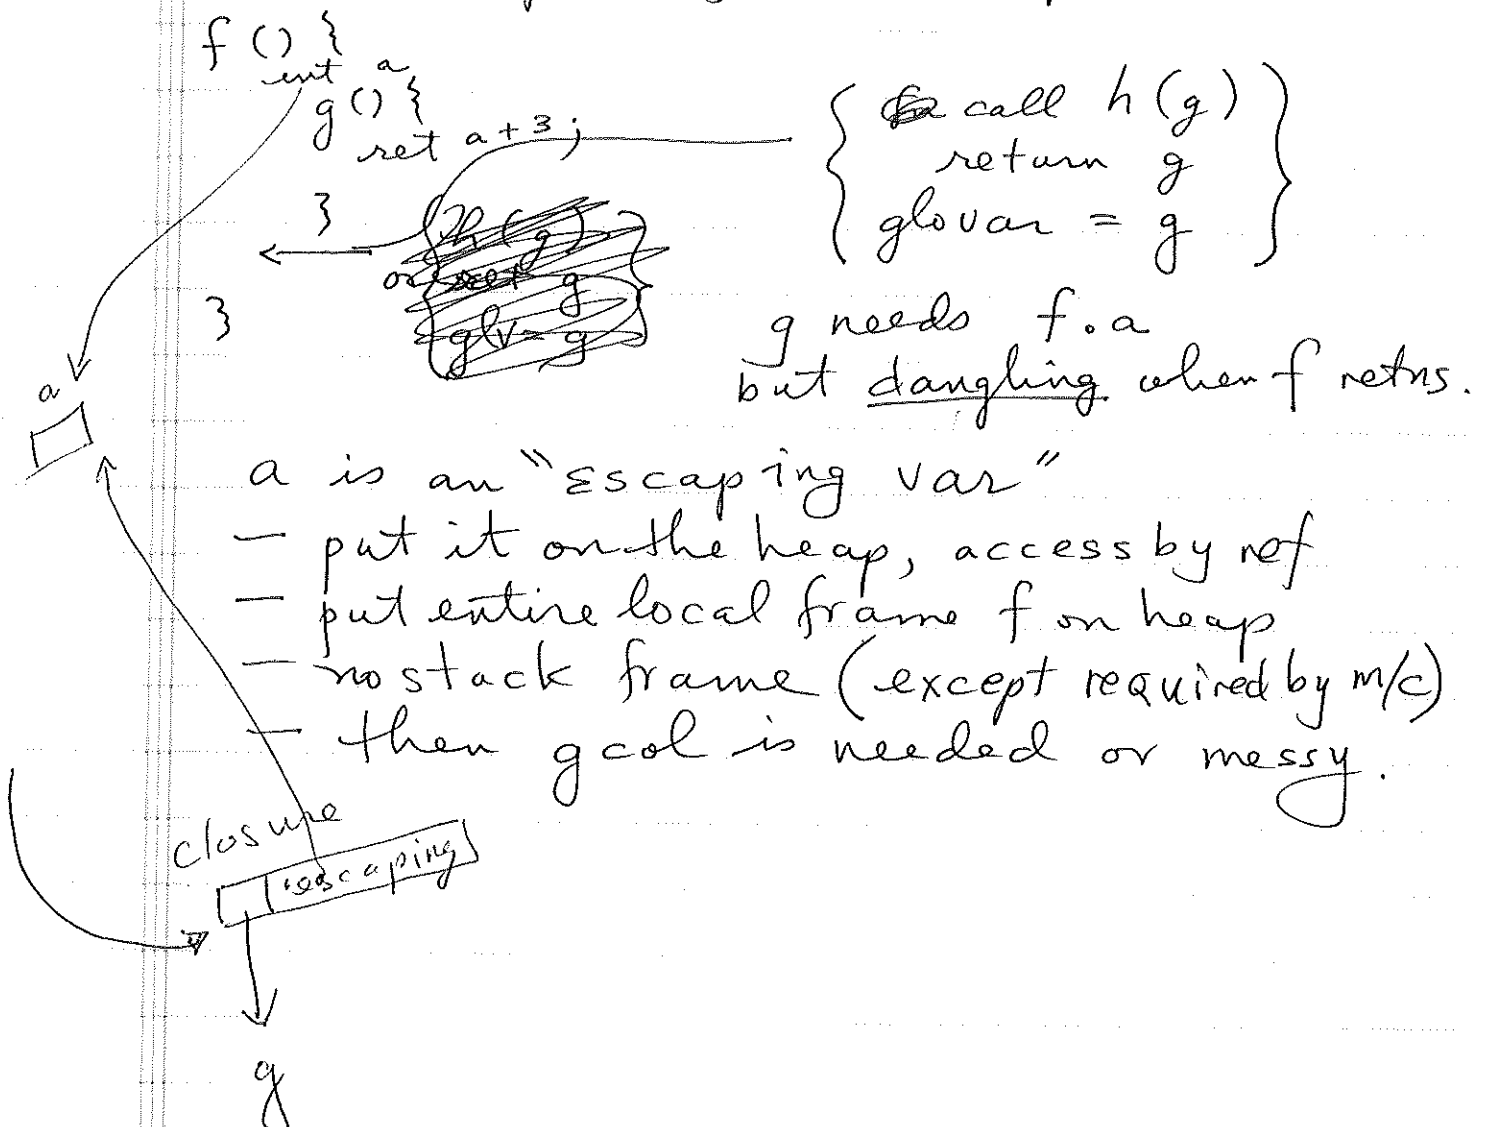
— no nested procs
⇒ C → FP & global.

— nested:

— if no fn params, not a problem
just use stack

— however: can't return local by ref
or: assign & local → global.

— what about returning a local fn,
or passing it thru param



Dynamic Scope

("local" in Perl)

Runtime

- deep access — search stack each access
- shallow access — ~~push new var on~~
 - each var is a stack
 - local() → push new var
 - exit scope → pop it
- PL/1 called these "controlled vars"

PIC — Position Independent Code

- shared libs mapped to shmem (RO)
- can't have any abs addrs
- don't need to relink progs
- need dynamic linker

→ ~~main~~ GOT = global offset table.
= offsets of external symbols

ex: $gp = GOT_off + 8$
call next, r31

next: $gp = gp + r31$

rel to instr
that uses it

PLT1:

save regs
call dynlinker
function ID

linker zaps it to

PLT1:

$g1 = \text{abs addr}(f)$
call (g1)

table in RW
mem

trampoline

~~GMP 104B~~
~~page 7~~
~~week 4~~

Saving & Loading Regs

- preserve across call
 - caller save
 - callee save
- notice which insns change regs too
 - eg. VAX movc5 ~~changes~~ R0..R5
 - SPARC call - ch % 0

SPARC save } does not save regs
Restore } just spins reg win

PREFERENCE	caller reg	callee reg	all reg.
callee save	1ST	3RD	5TH
caller save	2ND	4TH	6TH

↑
fewer
regs

less code

~~1ST callee save~~

1ST ee sv er reg. → need reg mask

- HW support

- w/o more efficient save allregs

2ND easy for retn inst

- global gotos difficult ... must find reload code.

- exceptions

3RD compiler knows regs in use - prolog must be backpatched

- maybe epilog has a backjump

4TH bitmaps only poss.

~~need~~ need reg mask.

VAX call is bitmap mask.

5TH easy if few regs - unrolled loop if few regs.

STMult.

6TH easy - should use utility routine

Heap Management

dyn storage allocation

- static + automatic (stack) not sufficient.



- static: fixed addr; saved call to call
- stack: - unneeded if \exists heap
 - used for efficiency
 - modern arch req it.

- heap: dyn stg. area (unrelated to heapsort)

- allocation (w explicit free)


- Fixed sized blocks

- maintain a list
- use links inside block for free list




- Variable sized blocks

- first fit or next fit
- need to keep extra pointers




- boundary tags. $\&$ malloc $\&$ byte ahead

~~free(p)~~ ~~free(int *p)~~



- becomes messy & bad for virtual mem
- frags the pages.
- split free blocks \rightarrow alloc
- coalesce when free.

malloc()



\leftarrow sbrk(): always mult of page size.

- compatible: new vs malloc
- portable: only use sbrk() syscall.
- localize. for other OS.
- minimize time
- locality \rightarrow paging & caching

Fastest free(){}: ~~malloc()~~

malloc(n) {
p = first free;

~~first free = (first free + n + 7) & ~7;~~

first free = (first free + n + 7) & ~7;

if (first free > break) {

break = (first free + extendant
+ pgsiz - 1) & ~ (pgsiz - 1);

sbrk(break)
assert(OK)

}
return p
}

~~p = end heap;~~
~~end heap += 5;~~
~~if (end heap > brk)~~

\leftarrow alignment.

\leftarrow magic bit put
0xF?
instead

pgsiz = 2^K
(1K to 16K?)

extendant
= m * pgsiz.

\uparrow
magic #

alignment: 8 for double.
(strictest hardware)
req.


2^K : 8 or 16?

\leftarrow pentium base reg
granularity?

actual alg

Boundary tags:

- size fields before and after chunk
- coalesce bordering unused chunks
- traverse all chunks fwd or bwd.



Bins: - maintain chunks in bins, by size

- 128 logarithmic size bins.
- smallest - first, best fit order.

size: 16, 24, 32 ... 8×64 9×64 10×64 ... 2^{31}
 exact bins sorted bins.

- use wilderness chunk only when all others fail (assume ∞ size)
- sbrk() - extends wilderness.
- mmap() - chunk between heap & stk
 - only for very large req.
 $ex \geq 1MB$.
 & not avail in exist.

syscalls are expensive

caching: lazy coalescing

- most pgms alloc/free chunks of a few fixed sizes.

ARENAS:

array of fixed sized chunk
 good for structs (classes)

less useful. ~~less~~ for arrays