Linking & Loading

Classical unix a.out format

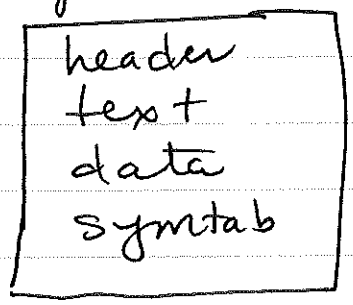

nm
strip

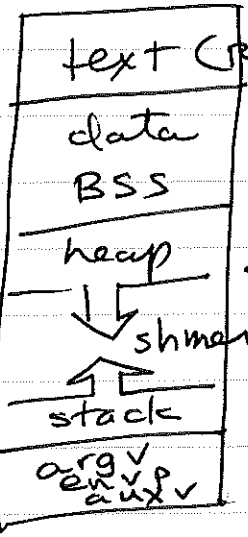

fork() - dups process mem image
data + BSS + stk = COW

vfork() - suspends parent until child exec a exit

execve() - relse addr space
- keeps open files
- alloc stk stg
- loads a.out
- jumps to entry address

execlp execl execle
↓ ↓ ↓
execvp → execv → execve

---

file id by magic #

a.out = assembler output (PDP11)

usual * ELF = executable & linkable fmt
COFF = common object file format
DWARF = debugging with attributed record fmt.
jar = java archive (0xcafebabe)
shell = #! path to interpreter

Link (2)

(ld) linker – combine obj files into exec image

(execve) loader – loads pgm from disk into mem ⇒ process.

relocation – changing addresses to move to other areas mem.
old – clunk & load
with VM – new (assign abs V. adr)

symbol resolution –
def one obj; undef others.

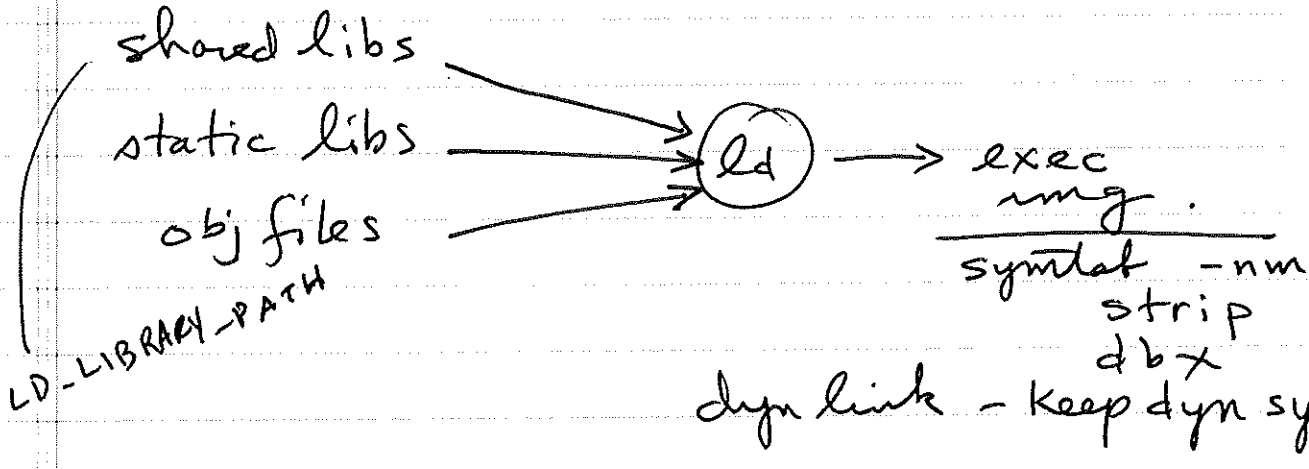

symtab -nm
strip
dbx

dyn link – keep dyn sym.

Reloc & code mod

– asm generates unreloc addrs (rel start of segment)
– also reloc dict

ex sethi %hi(foo), r1

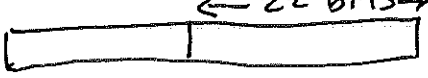

sym
foo → bits 21:0 = bits 31:10

CH-2

Architecture — Link(3).

need ABI spec.

ch 3 – machine interface

- data rep
- big/little endian / align req.
- registers, stack frame.
- call/return conventions
- dyn stack?

ch 5 – prog load & dyn link

API

struct layout

addresses: absolute
PC-relative
reg ± offset

virtual mem?
- if not loader must adjust addresses
- or must run in fixed loc.

mapped files
- exec faults in page from exec
- demand zero for BSS
- fault from swap

shared libs

PIC = position independent code.
- must run any location
- use relative jumps
-

CH - 3

Object Files

Link (4)

Header – magic #, list of segments & sizes.
Text – object code.
Data – RO data, RW data, ZW data (BSS)
relocation dict
global symbols (def & undef)
-g debugging info – local vars
source files & lines

Formats

a.out

text ← page boundary
data
BSS ← no boundary

header: magic
text siz
dat siz
bss siz
sym siz
entry adr
txt reloc size
dat reloc size

history

PDP-11
magic # 0407 (octal)
br pc+7words

– newly alloc mem auto init = 0
– ∴ arrays of ∅'s in bss.

int a[] = {0,0,0,0}
int a[4]; — better.

text pages ⟶ RO map

data pages ⟶ cow map.

relocatable a.out

- linkable files
- relocation table
  - entry: address - within module.
    - offset from start of module.
    - bit field spec.
    - is it external def?
    - PC rel?
    - name index
  - move module → adjust all addrs.
- symbol table
  - arb long strings
  - rep by symtab offset.

---

a.out → COFF → ELF

- Unix ELF
  - allows dyn linking
  - allows C++
  - debug fmt = DWARF

3 kinds:
- relocatable - created by compilers & as.
- executable - cre by linkers.
- shared = library files

ELF = set of logical sections
- each section ⇒ segment

ELF

Link (6)

extends T+D+B.

magic # [4] = "\177 ELF" = "~ELF"

177
01 111 111

reloc file = collection of sections
exec file = collection of segments.

sections: .text (PROG)
.data
.rodata = read only
.bss
.rel.text
.rel.data
.rel.rodata } reloc info for sections.
.init
.fini } startup & finishup code for C++ classes.
.symtab
.dynsym } symtabs
.strtab
.dynstr } strings
.got global offset
.plt proc linkage
.debug syms
.line = line #s to source.
.interp = path to #!

MIPS
short data, long data
short bss, long bss

pseudo
UNDEF
COMMON = common blocks
ABS = abs syms

ELF exec
= collection of segments for mmap
= similar set of names.
each seg is concat of input sections

mapped text seg = header + text + rodata
data seg = rw data
ext by BSS via demand zero
ext to heap by sbrk()

~~init & fini calls~~
init & fini in text seg
- called befr/after main

ordering can be critical
init - C++ static ctors
fini - C++ static dtors

ocaml - ordered by module
all global code.

CH-4

Storage Allocation

— static.

set of like sections → segment.

ex: A size 0x600
B size 0x4F
C size 0x701

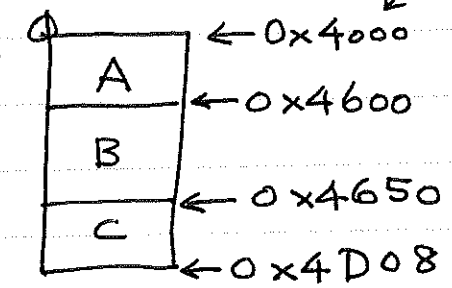

align each section to max hardware align
SPARC = 8
even if not needed
Pentium: align load is faster

not page bndy

reloc as concat them together.

Fortran COMMON blocks overlaid
- sized by largest.

Collect all text seg together
" " data " "
" " bss "

bss not in reloc or exec, but alloc demand zero.

Link(9)

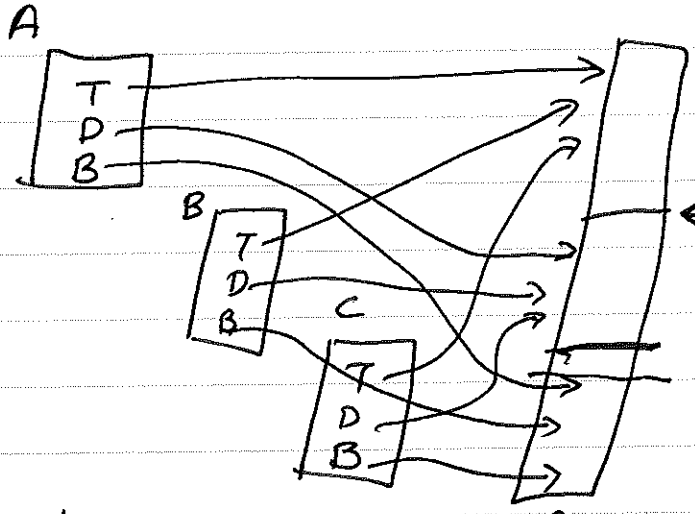

text seg round up to full page
data+bss start next page bndy
& round up. pg size.

Fortran IV — no dyn stg alloc.
- had to statically alloc all data
∴ tweak arrays & recompile
- share via common
- modules declare common arrays dim 1
- BLOCKDATA section actually alloc common.
just recompile this
COMMON → BSS

C++ duplicates

- virtual fn tables
- templates ⇒ macros

~~recompile template~~ ∀

∀ file ∀ class ⇒ compile VFT
∀ file ∀ type ⇒ compile template and VFT.

/SUNWspro/CC
→ uses subdir to cache compilations

Initializers & Finalizers

C: atexit(...)
C++: static vars' ctors/dtors.
- init seq = list of ptrs to startup.
- fini seq =

ordering problem in C++
ocaml → do ordering by linker order
can't have topological dependency cycle

CH-5 Symbol Mgmt Link (11)

Names
- global syms def
- global sym ref undef
- seg names
- local syms for debuggers
- line # info

symtab
- like compiler but flat
- name max len?
- store hash in table
  do strcmp only when hash matches

```
struct sym { char *name;
             int hashvalue;
             struct sym *link
           }
```

global sym table

each sym — def exactly once from one module
- ref *.

Link (12)

symbol resolution
- 2nd pass

refs - data ptr
- in insn
- multiple insns

ex: sethi %hi(X), g1
or g1, %lo(X), g1
22 bits, 10 bits.

Special syms

etext - end of text
edata - end of data
end - end of BSS
__CTOR_LIST__

Name Mangling

- avoid collisions: user ids vs gensyms

Cstyle:
old _main new main
gensym raw. •gensym

C++ encoding : type & scope

V vs C::V

f(int) vs f(float) open >>

nm

externals: foo → foo

func(float, int, unsign char)
→ func__FfiUc

classnames: #&name. ex: 4Pair

First::Second::Third → Q35First6Second5Third

cl::fn(void) → fn__2clFv

* → __ml

|= → __aor

⇒ link time type checking

Debugging info

- line # info.
- types, locations of vars
- each fun local var map.
- duplicate header files
  #include <iostream> X 20 copies

strip.

CH-6

Libraries

Link (14)

library files : archive formats
symtab → seq of obj files
- saves space on disk
- each fn sep obj

('.a') suffix
'ar' cmd
linkers make pass over libs.
-lm ⇒ /usr/lib/libm.a
/usr/lib/libm.so
option must be last.

we

CH=7 Relocation Link(15)

1st pass - lay out position of all sections in segment
arranges segment addresses

2nd pass - fix up storage addrs of all global syms.
data adrs = absolute
code = absolute | PC relative.

load (exec) relocation
- only if sys has no Virt mem.

- call X → needs rel diff X.
- handle partial addresses

sethi$_{22}$ lo$_{10}$ hi-ord 30 bits.

~~CH=8 Loadi~~

CH=8

Link(16)

Loading & Overlays

load :- read header
- allocate addr space
- mmap program
- zero out bss
- create stack.

overlays : used for old (non-VM) systems.
- let pager/swapper handle.

PIC = position independent code
- within module all code refs relative
- data at fixed offset from code

linker creates GOT = global offset table.

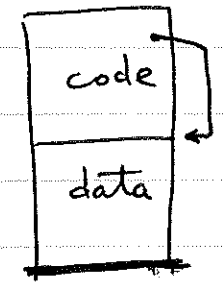


```
call  L
L: add  $GOT, ra
        ↑
      offset.
```

GOT then points at data

CH-9

Shared Libraries

Link (17)

- static link → put in exec
- shared → loaded @ exec time or dyn.
- exec → lib
- at load time
  startup code maps them into address space.
- static linked shared

- reserve address space.
  ex: VAX 0x00 — to 0x7F — user
  0x80 — to 0xFF — shared.

startup code
- finds libs & maps them into address space
- shared or COW
- may be in: OS executable dyn linker.

- must be avail when pgm run
- if changed?
- link time: sym bound to addresses
- can't change entry addresses.

shared static lib

- header
- init routine
- JUMP TABLE
- code
- global data
- private data

---

2 libs — shared lib itself
— stub lib linked into exec.

JUMP TABLE: read: jmp read'
write: jmp write'
⋮

change code doesn't affect jmp table
can't rearrange jmp table.

/usr/lib/libc.a

Malloc hack — your own version
~~you have~~ — but libs still call original

∴ extern void *(*mallocptr)(size_t);
#define malloc(s) ((*mallocptr)(s))

in lib init
#undef malloc
mallocptr = &malloc

stub in app, not lib, so: Link(19)
if app has malloc, it's used
else lib's malloc is used

Linking and Loading

CH = 10 Link(20)

Dynamic Linking & Loading

fetches libs from std /usr/lib
or compiled in -L libs
also: setenv LD_LIBRARY_PATH.

advantage: easier create than static shared
easier update
semantics more like unshared.

disadv: slower: link done every run.
larger - need symtab
dlibs must be backward compat

uses PIC
needs GOT = global offset table
PLT = procedure linkage table
- "lazy" evaluation

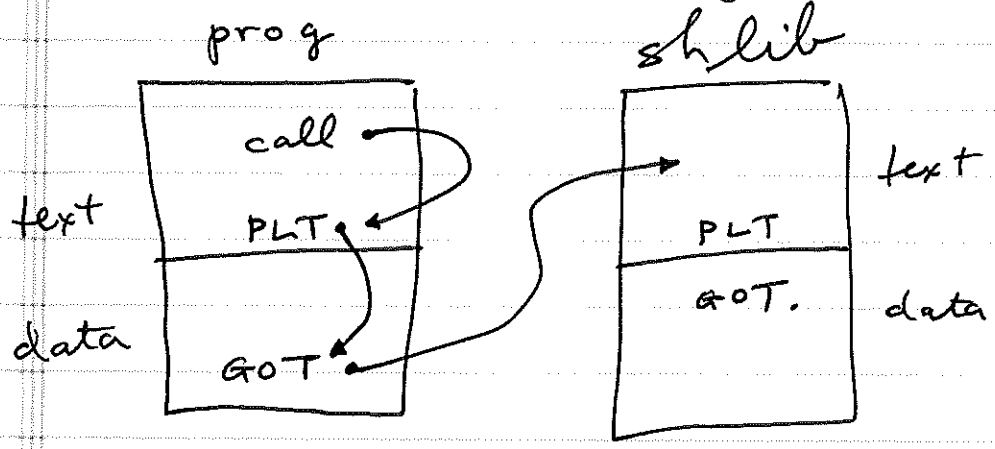

PLT

Link (21)

in exec	in mem
PLT0: unimp unimp unimp	PLT0: save %sp,-64,%sp call dynlink nop
PLT1: unimp unimp unimp ⋮	PLT1 .word identification unimp unimp
PLT101 sethi (.-PLT0),g1 ba,a PLT0 nop	PLT101: sethi (.-PLT0),g1 sethi %hi(printf),g1 jmpl g1+%lo(printf),g0
PLT102 sethi (.-PLT0),g1 ba,a PLT0 nop	PLT102: sethi (.-PLT0),g1 sethi %hi(fopen),g1 jmpl g1+%lo(fopen),g0

1. at load time, dyn linker alters PLT0 and PLT1 to xfer ctl to it "identific" word @ is a message to itself.
2. prog calls "printf", which is really "PLT101" sethi puts dist to PLT0 into $g1_{[31:10]}$
   $g1 = (101 * 12) \ll 10$
   branches to PLT0 calls dynlink with new frame
3. using ident, it finds data structs. $(g1 \gg 10)/12$ gives index where "printf" info is

4. dynlink alloc space for "printf" text & data, loads it, unwinds stack, fixes PLT 101, jumps to "printf"

FIX NOTE: fix PLT 101
- must update re-entrant in case interrupt/signal
atomic update: 3rd word.

① ⇒ sethi
ba,a
nop → jmpl } annuls delay insn

② ⇒ sethi
sethi
jmpl } 1st sethi irrelevant
jmpl delay slot uses next PLT sethi
↳ OK because jmpl addr computed

/usr/lib/ld.so.1
/usr/lib/libc.so.1

libs